



 Cached During Inference


Multi-Head Attention (MHA)



Grouped-Query Attention (GQA)



Multi-Query Attention (MQA)



Multi-Head Latent Attention (MLA)

