



Baseline: Prompted/RLed agents



Learns action selection *implicitly* from internal confidence (**Prompted**) or training statistics (**RL**)

Calibrate-Then-Act

First, estimate priors



CTA: Enables *explicitly* reasoning over prior probabilities in the decision making process

WHAT CTA ADDS TO THE PROMPT

Additional context:

- Estimated format likelihoods are provided below.

Format likelihoods:

{prior}

What the model then reads out of it:

The delimiter is most likely to be ' ; ' with a probability of ~0.85 ... But I'm not 100% sure, so maybe I should run some unit tests to confirm.